

Code of Conduct

Fair and Responsible Use of AI in the AIOps Coursework

Welcome to the AIOps coursework. In modern engineering environments, AI-assisted coding tools (such as GitHub Copilot, Gemini, and Claude) are standard instruments for accelerating development. We encourage you to use these tools to enhance your learning and productivity.

However, the core objective of this course is for **you** to build a deep, functional understanding of machine learning infrastructure, model deployment, and lifecycle management. Relying on AI to bypass the cognitive work of systems design defeats the purpose of your education.

To ensure academic integrity and align with the principles of responsible AI, all students must follow these guidelines for AI use in assignments and projects.

1. Core Principles

- **Accountability:** You are the principal engineer of your submissions. You are strictly responsible for the accuracy, security, and performance of any code you submit. If an AI generates a flawed Apache Airflow DAG or an insecure Docker configuration, the responsibility falls entirely on you.
- **Transparency:** Document and declare the use of AI generation. Engineering is about reproducible and transparent processes.
- **Comprehension:** You must be able to explain, modify, and troubleshoot every line of code, infrastructure configuration, and deployment script you submit, regardless of its origin.

2. Acceptable Use (The "Dos")

You are encouraged to use AI tools as a learning assistant and pair-programmer for the following:

- **Boilerplate Generation:** Generating standard configuration files (e.g., basic `docker-compose.yml` structures, DVC initialization scripts, or MLflow tracking boilerplate).
- **Debugging and Troubleshooting:** Pasting error logs (e.g., from a failed Ray cluster deployment or a Prometheus metric scraping error) to understand the root cause and identify fixes.
- **Syntax Translation:** Converting working logic from one framework to another (e.g., translating a local Python script into a PySpark job).
- **Conceptual Clarification:** Asking for explanations of complex MLOps concepts, architectural patterns, or API documentation.

3. Unacceptable Use (The "Don'ts")

The following actions constitute academic misconduct and a violation of course policy:

- **"Zero-Shot" Homework:** Copying and pasting the exact assignment prompt into an LLM and submitting the unmodified output as your own work.
- **Blind Deployment:** Executing or submitting AI-generated orchestration scripts, CI/CD pipelines, or container setups (like Docker Swarm configs) without carefully reviewing them for security flaws or logical errors.
- **Falsifying Evaluations:** Using AI to fabricate model evaluation metrics, logging data, or simulated Grafana dashboard outputs instead of running the actual data pipelines and experiments.
- **Claiming Authorship of Architecture:** Using an AI to design your entire end-to-end MLOps architecture without engaging in the design trade-off process and technical decision-making yourself.

4. Disclosure and Documentation Requirements

For every assignment or project submitted, you must include a brief **AI_DISCLOSURE.md** file (or a dedicated section in your **README.md**). This disclosure must include:

1. **Which tools were used:** (e.g., GitHub Copilot, Gemini).
2. **How they were used:** (e.g., "Used Copilot to autocomplete MLflow logging syntax," or "Used an LLM to debug a Docker networking issue").
3. **Impact:** A brief sentence on how the tool shaped the final submission.

Note: If no AI tools were used, explicitly state: "No AI tools were used in the completion of this assignment."

5. The "Whiteboard Test"

Course instructors reserve the right to conduct a "whiteboard test" (verbal or written defense) of any submitted code. If you cannot explain the mechanics of your orchestration pipeline, how your data versioning works, or the logic behind your model tracking setup, it will be assumed the work is not your own, regardless of the AI disclosure provided.

By building a habit of transparency and accountability now, you are preparing yourself to build, deploy, and govern responsible AI systems in the real world.